




# A Workshop On Containerization

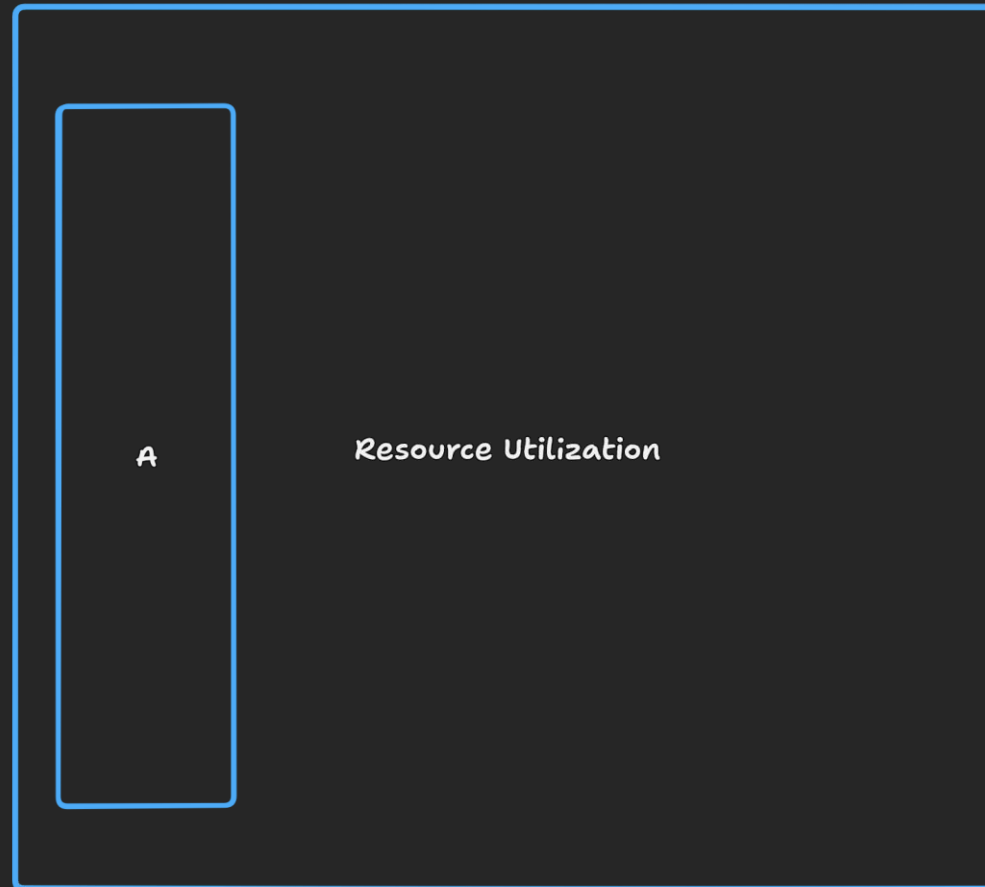


~~Why should I care?~~  
What are the benefits?

# Containerization Dominates



# An example




CPU  
16  
Cores

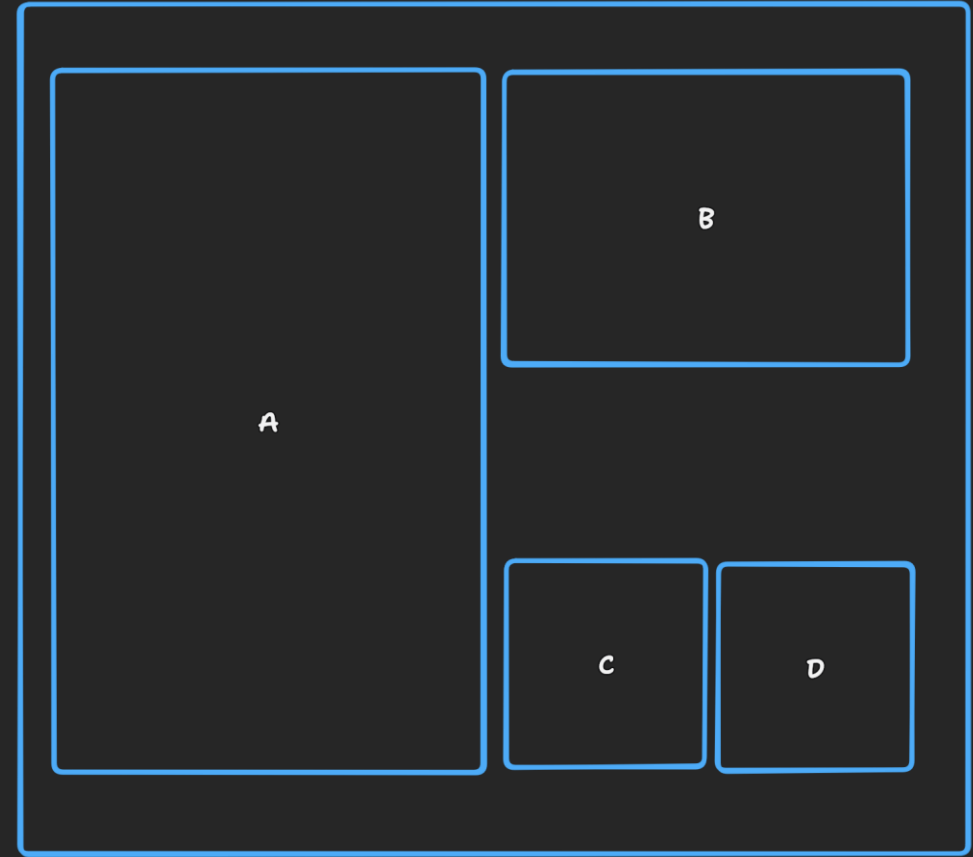
Memory  
64 GB

\$\$\$\$\$

# Sharing Model




Expectation



Reality

# Failures



- Dependencies were missing
- Wrong version of dependencies
- Wrong kernel API call
- Clashing ENVs



**IT WORKS ON MY MACHINE**




**THEN WE'LL SHIP YOUR MACHINE**




**AND THAT IS HOW DOCKER WAS BORN**

# Scale up





# Crashes




**Isolation**

**Set Resource Bounds**

**Portability**

**Reproducible  
aka**

**Write once, run anywhere**



**HOW DO WE SOLVE THIS?**



**CONTAINERIZATION, RIGHT?**




imgflip.com



**CONTAINERIZATION, RIGHT?**

# Virtualization

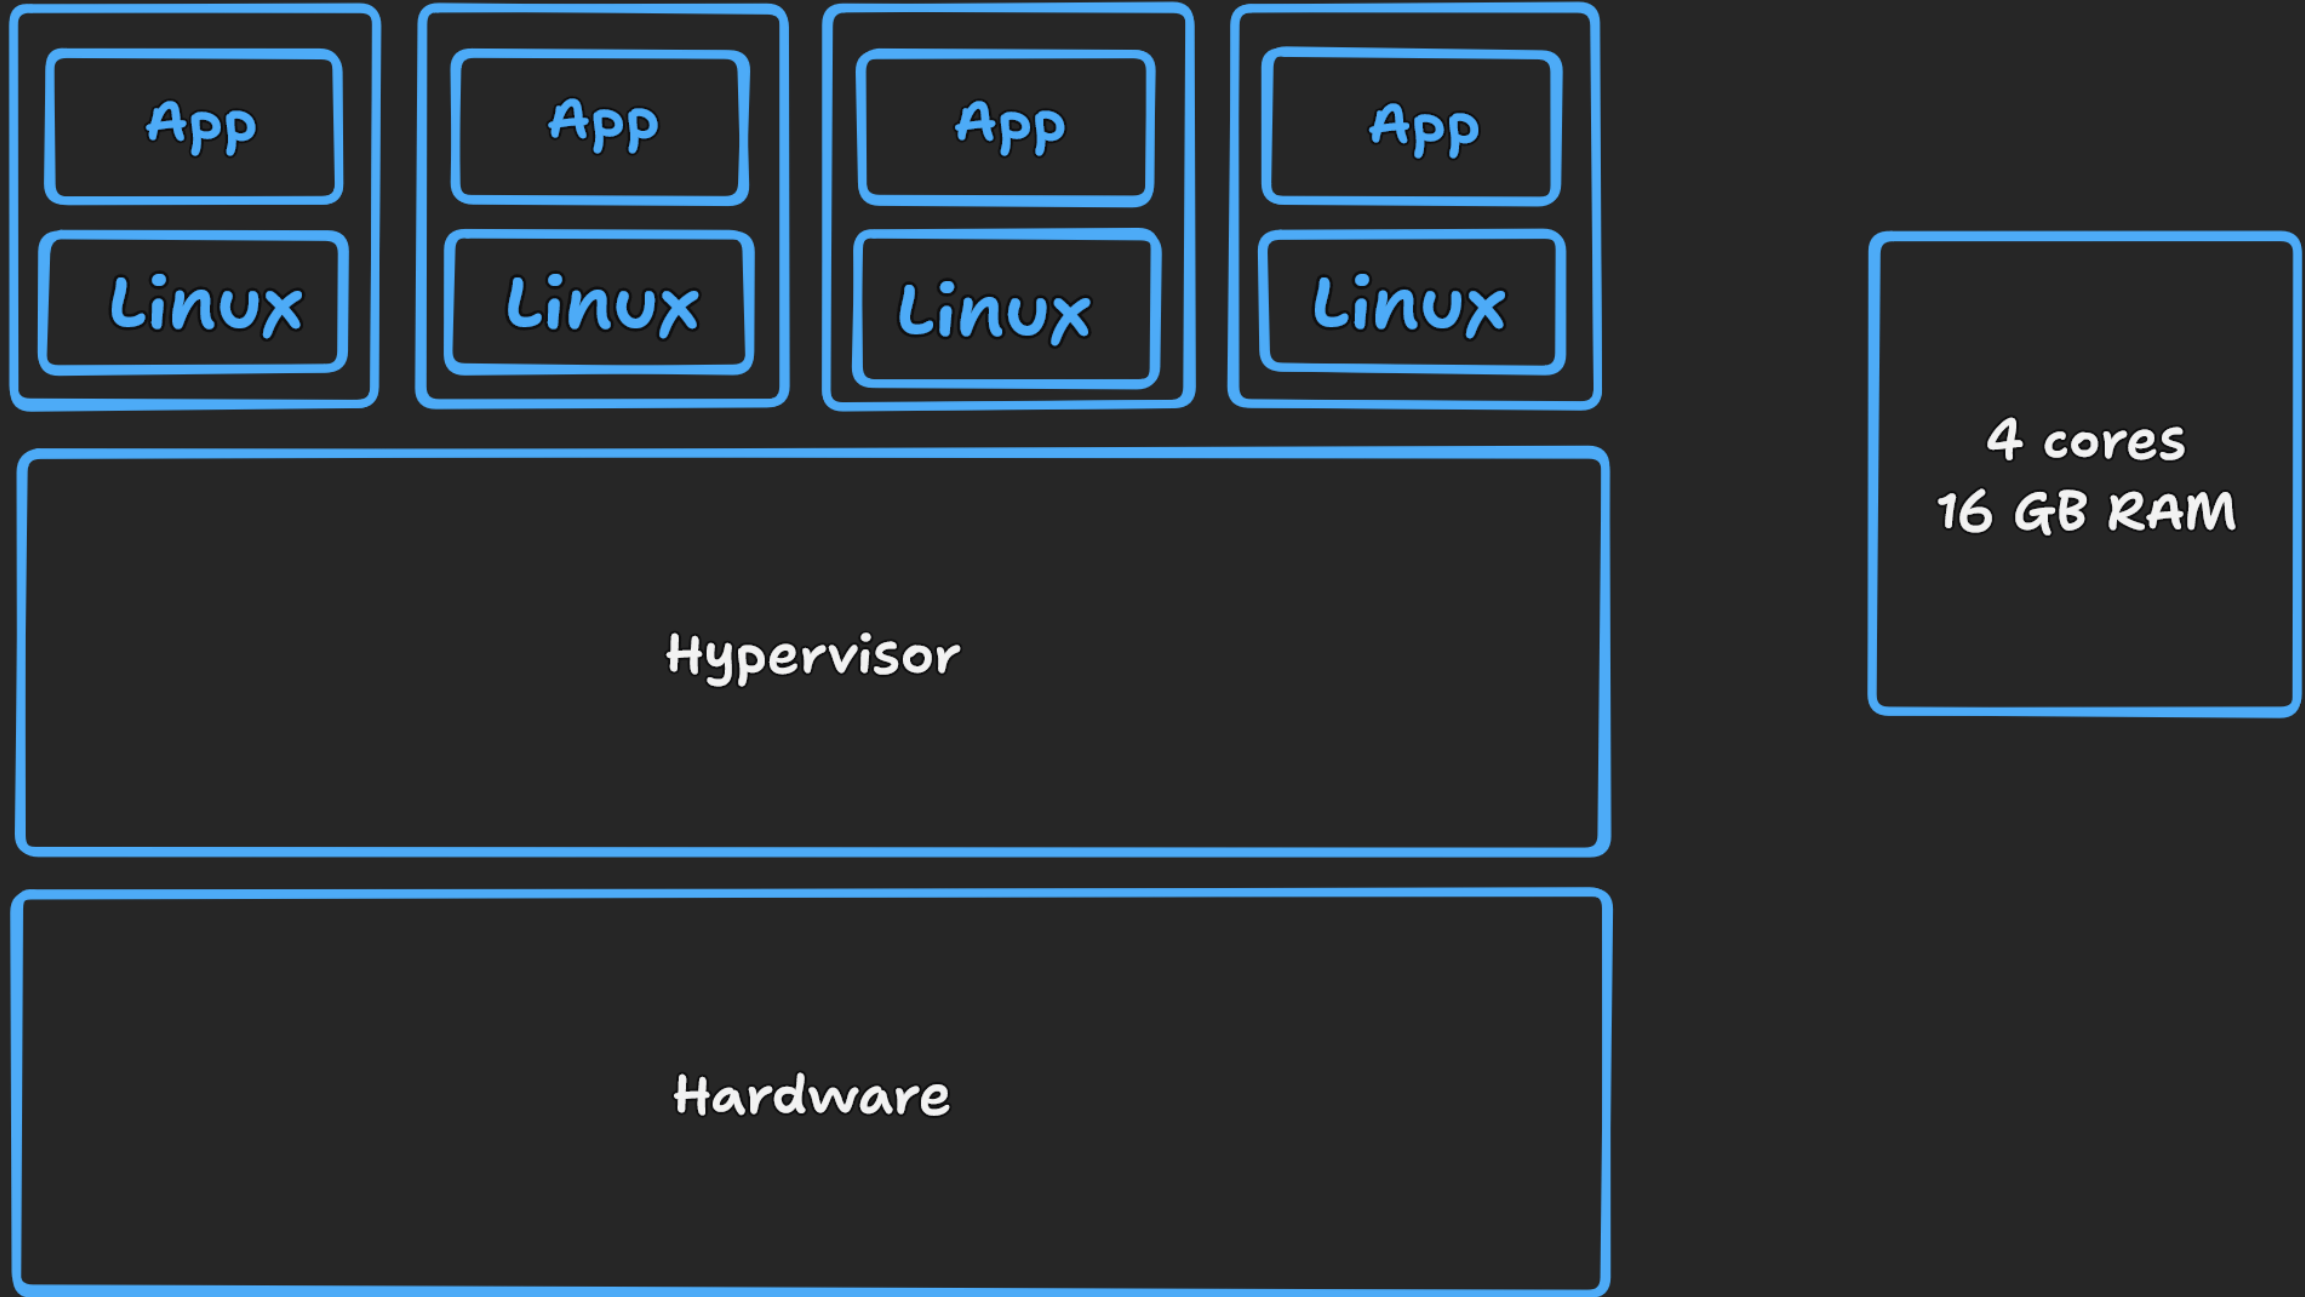




DigitalOcean

## Pros

- Strong Isolation
- Strong Security
- Scalable
- Flexible
- Portability
- Resource Optimization
- Cost savings
- Do more with less



App

App

App

App

Linux

Linux

Linux

Linux

Hypervisor

Hardware

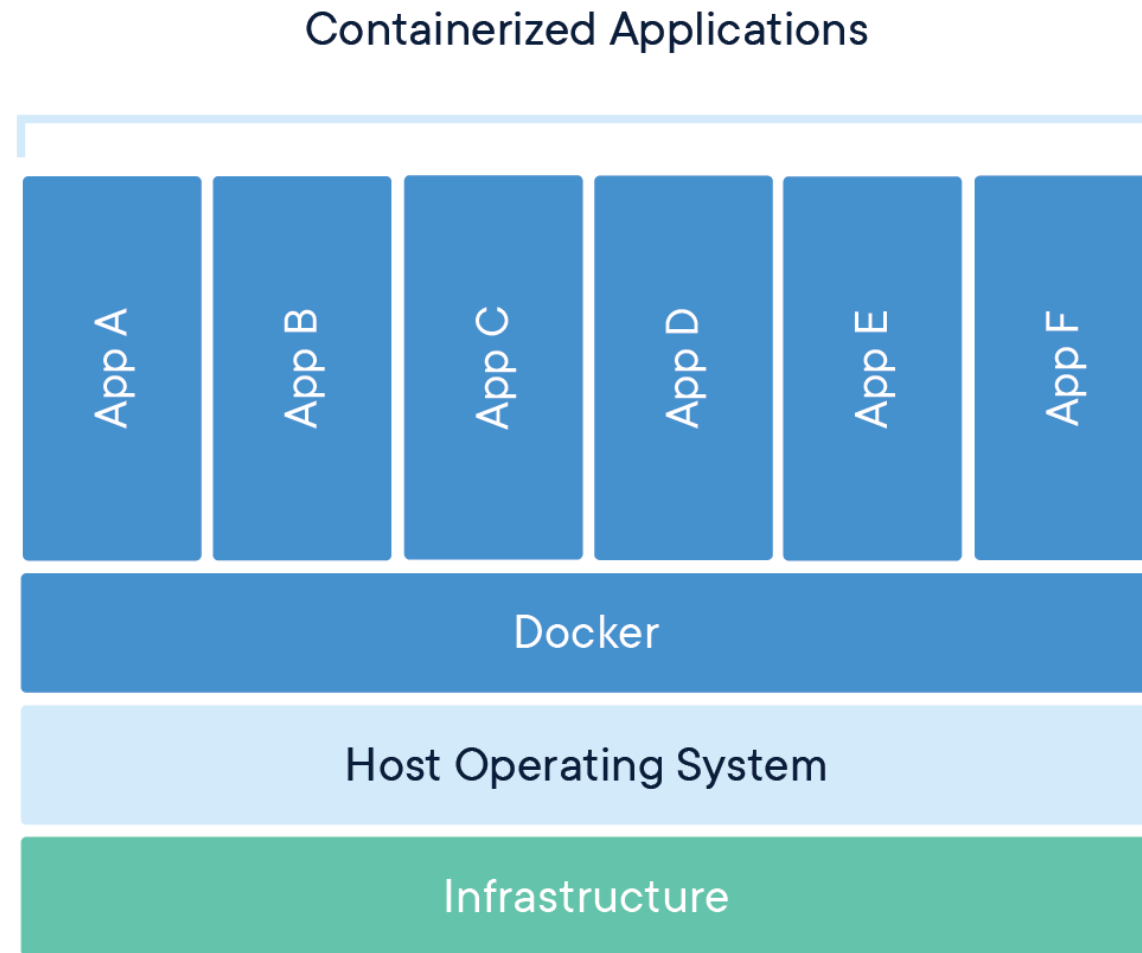
4 cores  
16 GB RAM



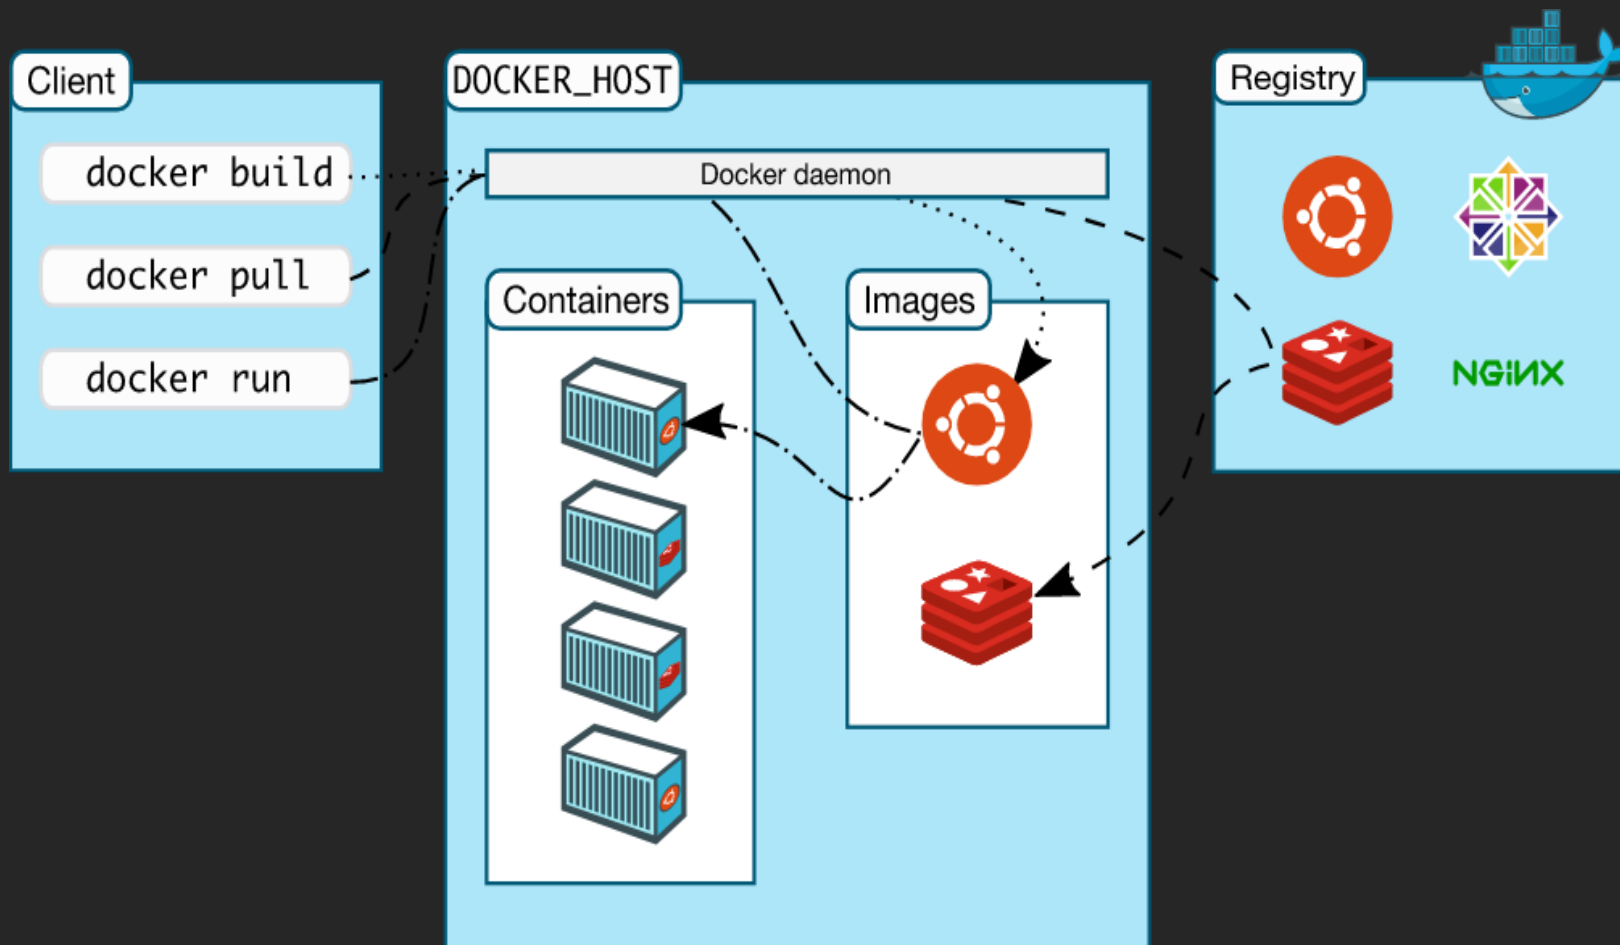
## Cons

- High Resource Overhead
- Worse Performance compared to native
- Slow to boot
- Limited Portability
- Cannot Rapidly Scale

# Docker



# Key Components



Daemon  
Image  
Container  
Registry

# Docker run

Creates a container from an image and starts it



```
docker run hello-world
```

# Docker image

## Get image information



```
docker image ls
```



```
cat python_docker_image | docker image import - python:latest
```



```
cat postgres_docker_image | docker image import - postgres:latest
```

# Docker pull



```
$ docker pull ghost
```

```
Using default tag: latest
```

```
latest: Pulling from library/ghost
```

```
Digest: sha256:d58cd8865658f66687d98c1fcbd228878aca9251431e6879694b31f2861d601a
```

# Example



```
docker run -d --name some-ghost \  
  -e NODE_ENV=development \  
  -e url=http://localhost:3001 \  
  -p 3001:2368 \  
ghost
```



# Docker container list



```
docker ps
```

# Docker inspect



```
docker inspect <container_id>
```

# Docker container



```
docker container stop <container_id>
```



```
docker container start <container_id>
```



```
docker rm <container_id>
```



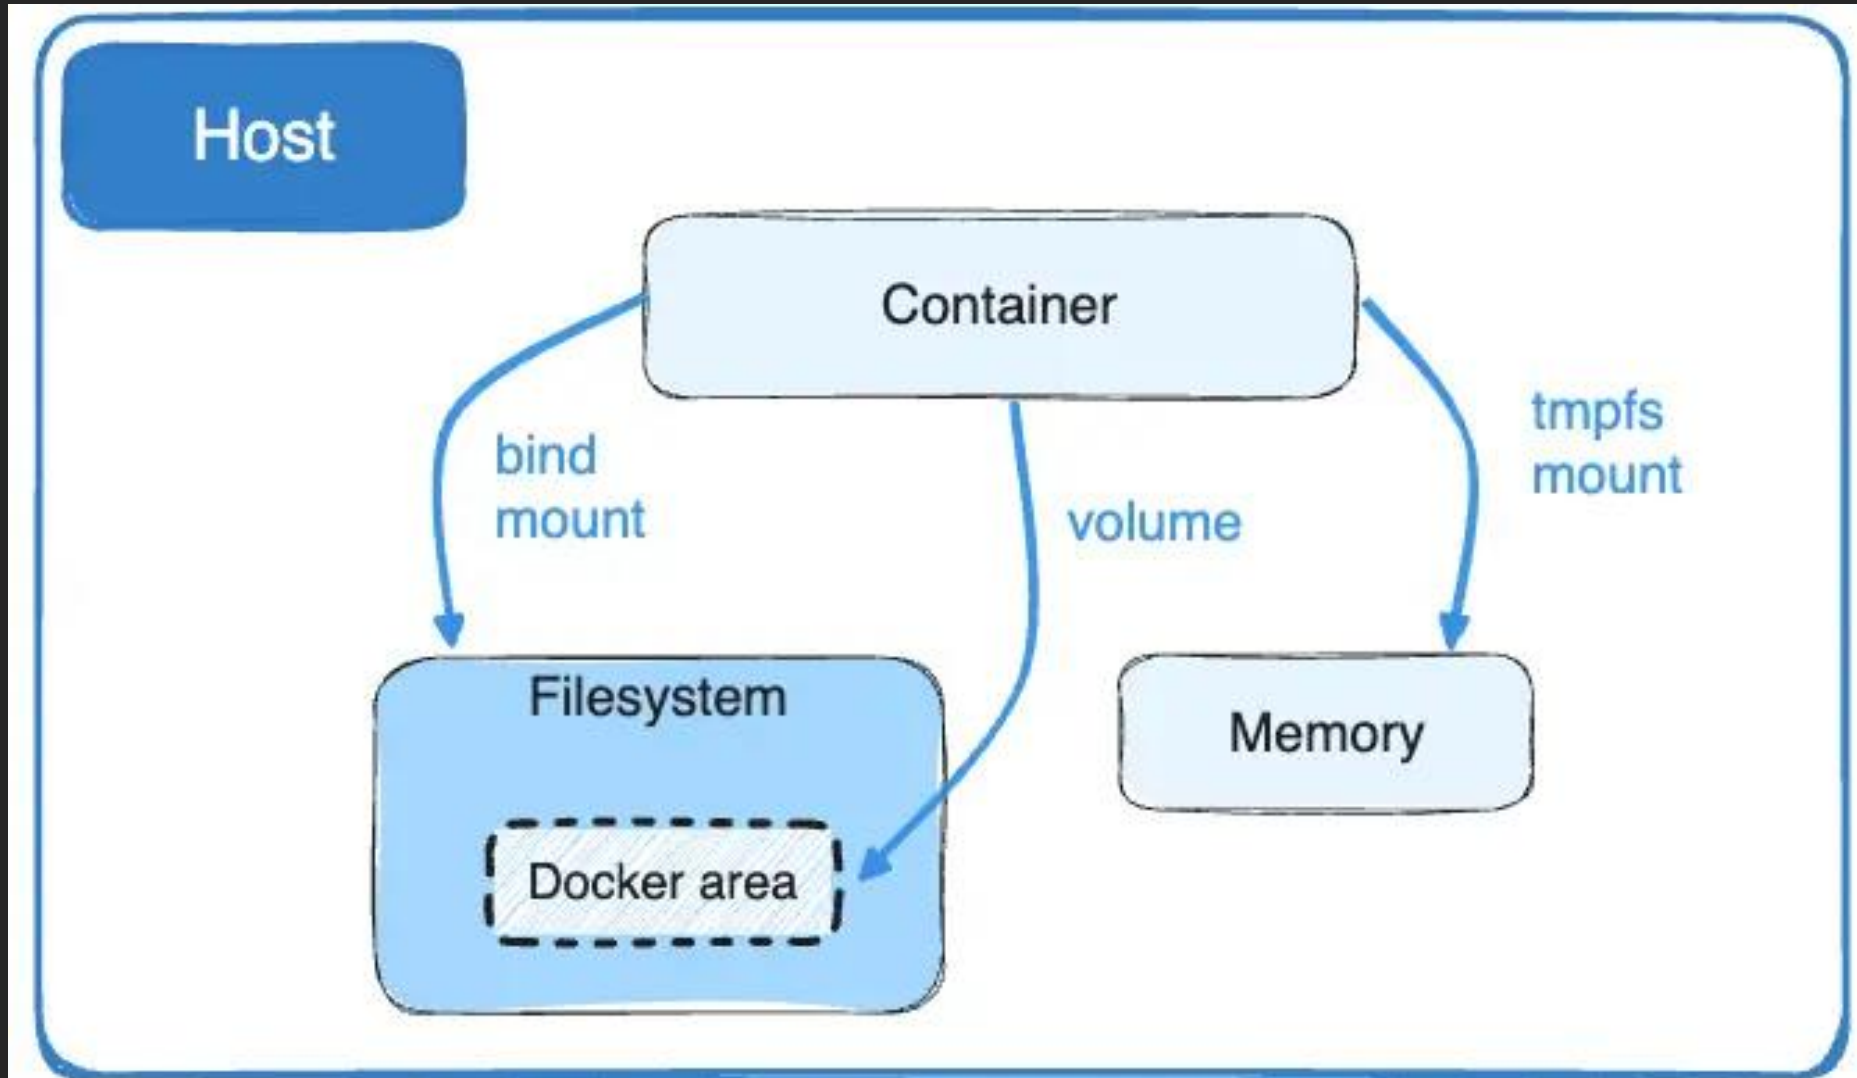
```
docker container attach <container_id>
```

# Connect to a container



```
docker run -it --entrypoint /bin/bash python
```

# Persist Data



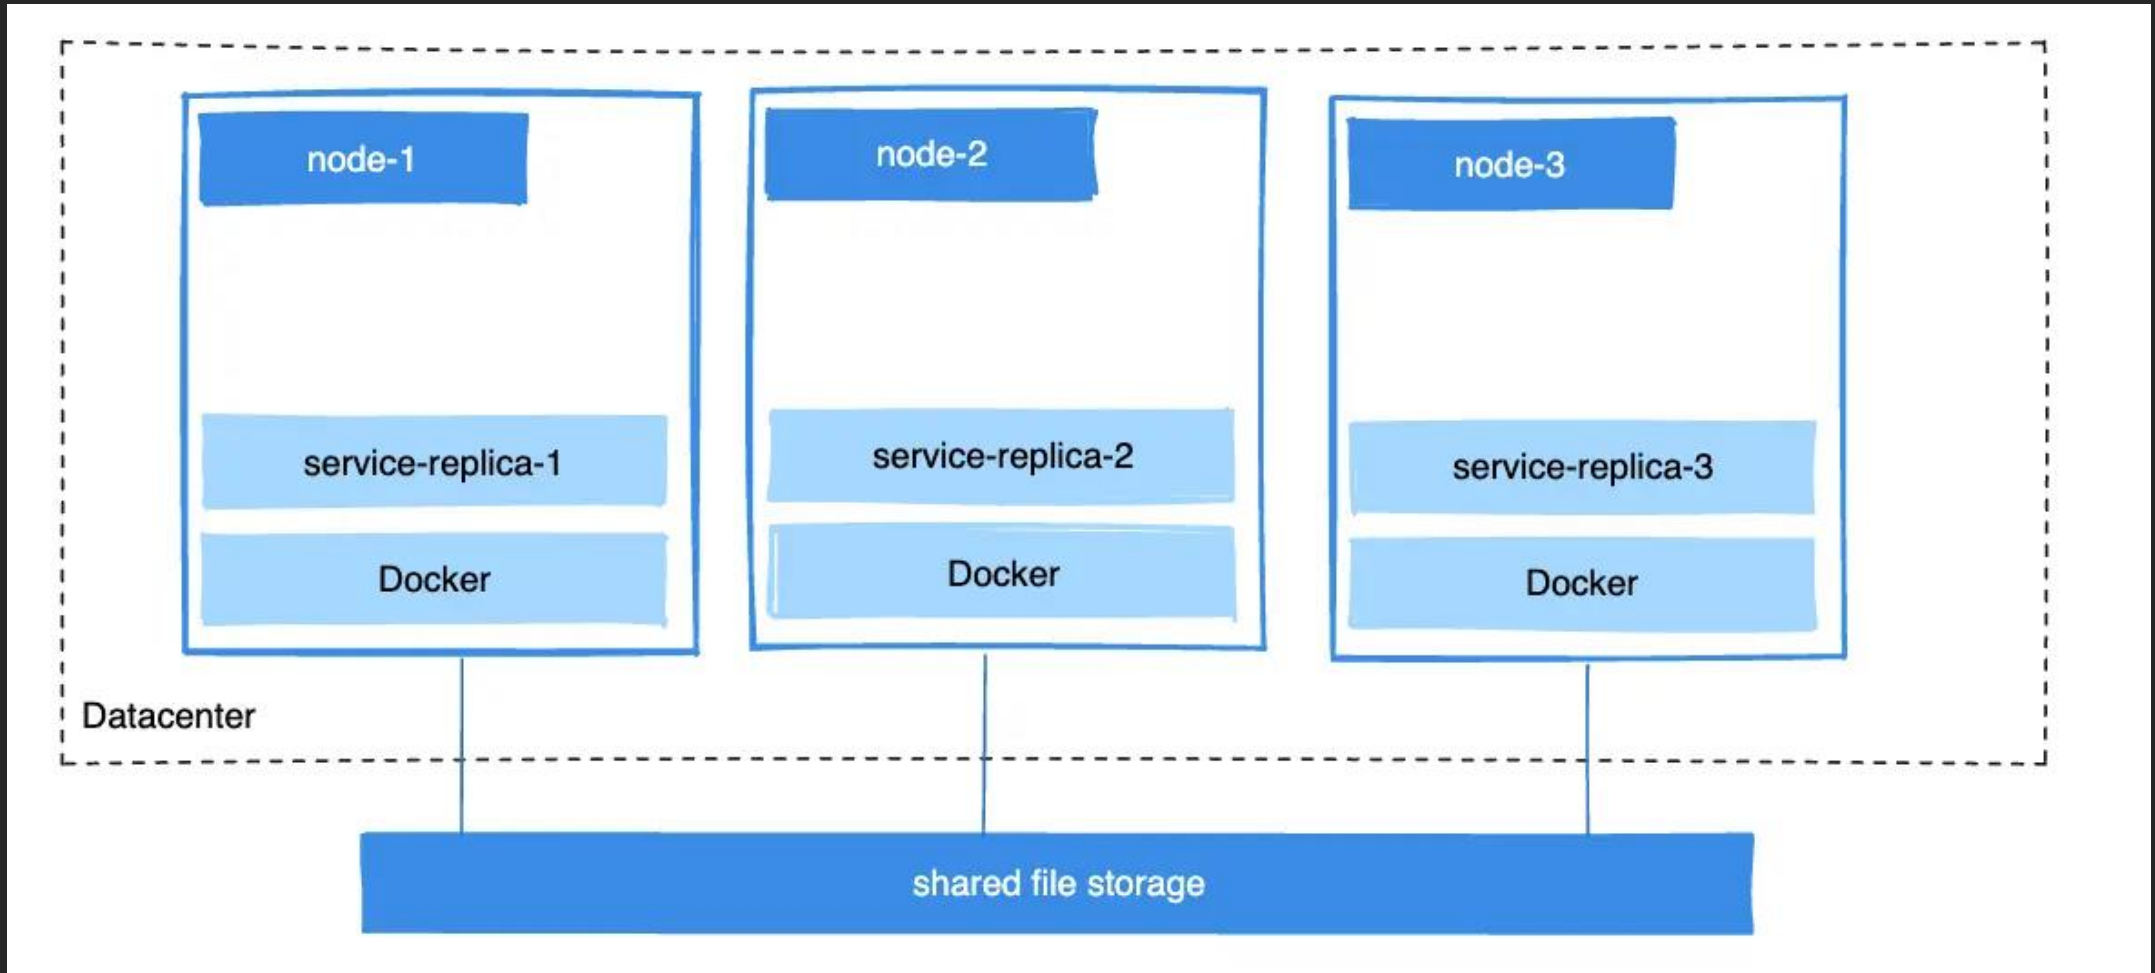
# bind mount



```
docker run -d -v <local_path>:/home/project my-python-image:latest
```




# Docker volume



Bind Mount	Docker volume
Dependent on the host OS file system	Completely managed by docker
Harder to backup	Easier to backup
Worse performance on windows and MacOS	Higher performance on windows and macOS
Cannot be stored on a remote host	store volumes on remote hosts or cloud providers, encrypt the contents of volumes, or add other functionality
Cannot be pre-populated by the container	New volumes can have their content pre-populated by a container


# Creating our own containers



Docker file



Docker Image



Docker Container

# Your first Dockerfile



```
FROM python:latest
```

```
RUN mkdir -p /home/project
```

```
COPY . /home/project
```

```
RUN pip install pandas matplotlib
```

```
WORKDIR /home/project
```

```
CMD ["python", "/home/project/main.py"]
```

# Docker build



```
docker build -t my-python-image .
```

# Dockerfile syntax

Command	Description
<a href="#"><u>ADD</u></a>	Add local or remote files and directories.
<a href="#"><u>ARG</u></a>	Use build-time variables.
<a href="#"><u>CMD</u></a>	Specify default commands.
<a href="#"><u>COPY</u></a>	Copy files and directories.
<a href="#"><u>ENTRYPOINT</u></a>	Specify default executable.
<a href="#"><u>ENV</u></a>	Set environment variables.
<a href="#"><u>EXPOSE</u></a>	Describe which ports your application is listening on.
<a href="#"><u>FROM</u></a>	Create a new build stage from a base image.

Command	Description
<a href="#">HEALTHCHECK</a>	Check a container's health on startup.
<a href="#">LABEL</a>	Add metadata to an image.
<a href="#">MAINTAINER</a>	Specify the author of an image.
<a href="#">ONBUILD</a>	Specify instructions for when the image is used in a build.
<a href="#">RUN</a>	Execute build commands.
<a href="#">SHELL</a>	Set the default shell of an image.
<a href="#">STOPSIGNAL</a>	Specify the system call signal for exiting a container.
<a href="#">USER</a>	Set user and group ID.
<a href="#">VOLUME</a>	Create volume mounts.
<a href="#">WORKDIR</a>	Change working directory.

# Docker network

## Communicate with Containers

## Communicate outside



# Types

Type	Description
Bridge	The default network driver. Bridge networks let containers talk to each other on the same host
Host	Removes network isolation from between container and the host machine
None	Completely isolate the container



```
docker network create -d bridge my-bridge-network
```




```
docker run -itd --network=my-bridge-network python
```

# Docker compose


```
docker-compose.yml

services:
  app:
    image: node:18-alpine
    command: sh -c 'yarn install && yarn run dev'
    ports:
      - 127.0.0.1:3000:3000
    working_dir: /app
    volumes:
      - ./app
    environment:
      MYSQL_HOST: mysql
      MYSQL_USER: root
      MYSQL_PASSWORD: secret
      MYSQL_DB: todos
  mysql:
    image: mysql:8.0
    volumes:
      - todo-mysql-data:/var/lib/mysql
    environment:
      MYSQL_ROOT_PASSWORD: secret
      MYSQL_DATABASE: todos
volumes:
  todo-mysql-data:
```

Config YAML



Docker Compose





```
FROM postgres:12-alpine
```

```
COPY ./db-init.sql /docker-entrypoint-initdb.d/db_init.sql
```

```
COPY ./data.csv /tmp/data.csv
```

```
CMD ["postgres", "-c", "log_statement=all"]
```



```
version: "3.4"
```

```
services:
```

```
  postgres:
```

```
    build: ./postgres/
```

```
    container_name: postgres
```

```
    ports:
```

```
      - "5432:5432"
```

```
    environment:
```

```
      POSTGRES_PASSWORD: "docker"
```

```
      POSTGRES_DB: "graph"
```

```
    restart: on-failure
```

```
    networks:
```

```
      - db-data
```

```
  app:
```

```
    build: .
```

```
    depends_on:
```

```
      - postgres
```

```
    container_name: my-app
```

```
    environment:
```

```
      DB_HOST: "postgres"
```

```
      DB_USER: "postgres"
```

```
      DB_PASSWORD: "docker"
```

```
      DB_NAME: "graph"
```

```
    networks:
```

```
      - db-data
```

```
networks:
```

```
  db-data:
```

# Command



```
docker compose up
```



```
docker compose up -d
```



```
docker compose down
```



```
docker compose stop
```



```
docker compose start
```

# What is a container?

A Standard Unit of Software

An Isolated Process



# “Magic” behind containers

- Chroot
- Linux namespaces
- cgroups
- LXC, containerd

# What is an image?

A collection of changes to a file system, affecting -

- Files
- Binaries
- Libraries
- configurations

# Principles of an image

- Images are immutable
- Images are composed of layers
- Each layer represents fs change


# Virtualization vs Containerization

Aspect	Virtualization	Containerization
Resource Usage	High overhead due to multiple OS instances	Low overhead; shares host OS kernel
Startup Time	Slower (minutes)	Faster (milliseconds)
Isolation Level	Strong isolation between VMs	Weaker isolation; shares OS kernel
Portability	Limited portability across platforms	Highly portable across different systems
Use Cases	Legacy apps, multi-tenant environments	Microservices, cloud-native applications

# Docker Alternatives



containerd



# Learn More

<https://medium.com/@saschagrunert/demystifying-containers-part-i-kernel-space-2c53d6979504>

<https://medium.com/@saschagrunert/demystifying-containers-part-ii-container-runtimes-e363aa378f25>

<https://medium.com/@saschagrunert/demystifying-containers-part-iii-container-images-244865de6fef>

